



# Operating Systems Security



# Learning Objectives

---

- Understand about the security threats in operating systems
- Learn about ways to protect

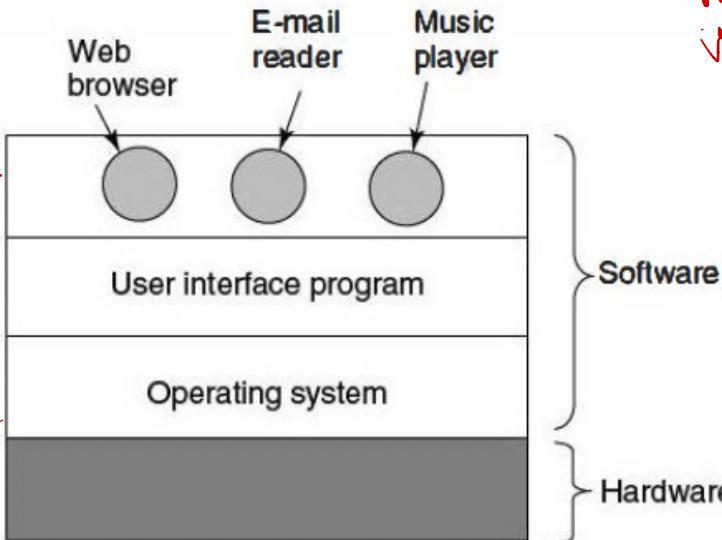


# What is an operating system?

- A layer of software
  - Provide user programmes with a simpler model of a computer
  - Handle the managing of resources

User mode must make requests to the kernel via system calls

restricted access to user resources  
can't access hardware resources directly  
device driver operations  
access hardware directly  
Kernel has complete control



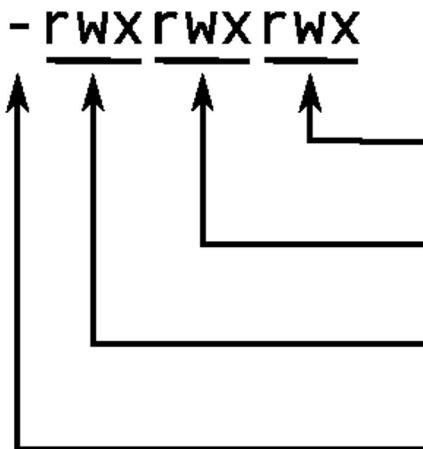
[source: 1]



# Why do we need protection?

- Confidentiality
  - Large amount of data is contained in systems
- Security issues
  - Human and nonhuman

↙  
*Attackers*



Read, write, and execute permissions for all other users. [source: linuxcommand.org]

Read, write, and execute permissions for the group owner of the file.

Read, write, and execute permissions for the file owner.

File type:  
- indicates regular file  
d indicates directory

↗ Malicious  
Software

↗ light processes

# Threats

- Exposure of data
  - Threats data confidentiality
- Tampering of data
  - Threats data integrity
- Denial of service
  - Threats system's availability
- System infected by viruses
  - Threats the goal of excluding outsiders

Threatens legitimate users

# Cryptography

---

- Cryptography may be used to ensure confidentiality and integrity
  - Symmetric cryptography
  - Asymmetric cryptography
  - Hash functions
- What if the keys are compromised?  
*Power lies in the keys*

# Trusted platform module (TPM)

→ stores keys

- Cryptoprocessor with some non-volatile storage for the keys
- Can perform cryptographic operations in main memory
- Can verify digital signatures
- Since implemented in hardware, it's fast too



# Trusted systems

- Can you build a secure computer system?
- Simplicity vs features *Tradeoffs*

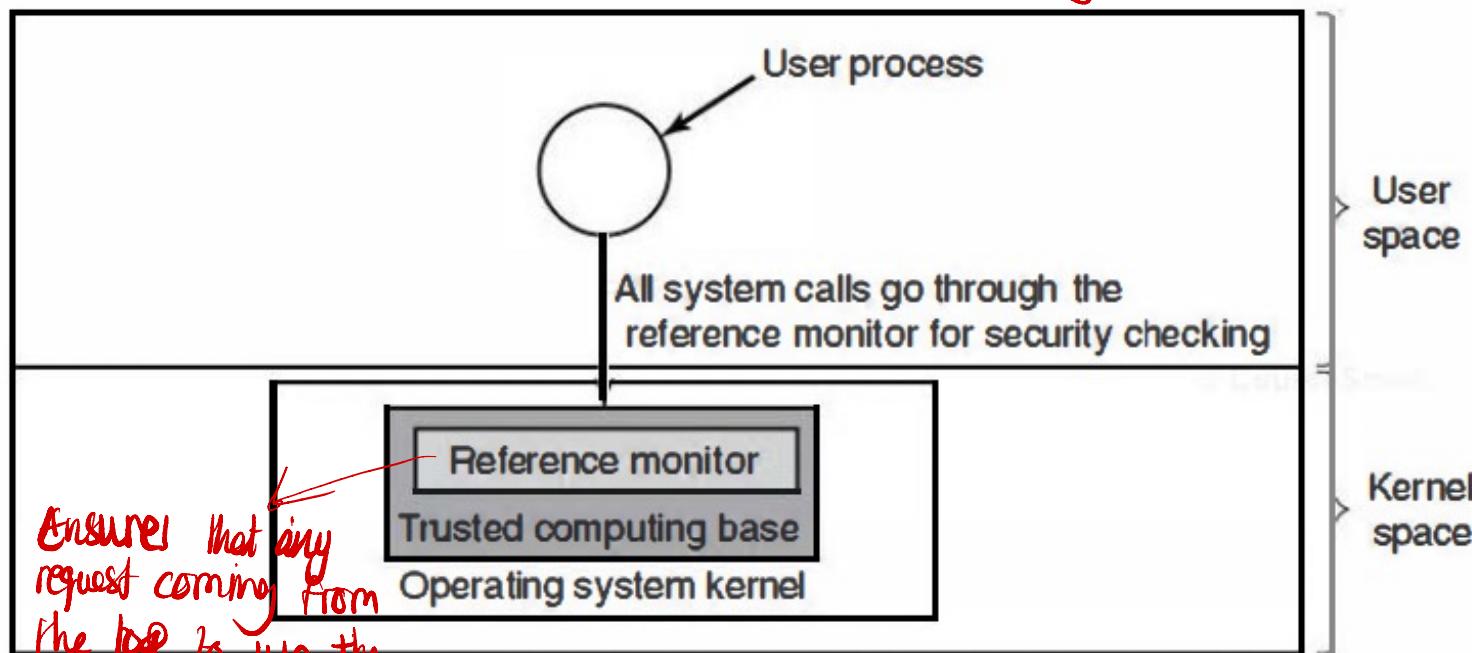
A screenshot of a Microsoft Word document window titled "Document3". The ribbon menu at the top includes Home, Insert, Draw, Design, Layout, References, Mailings, Review, Share, Editing, and Comments. The Home tab is selected. The toolbar below the ribbon contains icons for Paste, Cut, Copy, Paste Special, Undo, Redo, Find, Replace, and other editing tools. The main content area is a blank white page with a vertical ruler on the left side ranging from 1 to 6 inches.



# Trusted computing base

Recommended that code in the trusted computing base be kept to a minimal for security purposes

- Combination of hardware and software for enforcing security rules, *protects system from unauthorised access/tampering*



[source: 1]

# What else do we need?

## Models to restrict access

- Access matrix
- Multilevel security
  - Bell-La Padula model
  - Biba model

→ Breaks down to different security requirements (one solution or another)

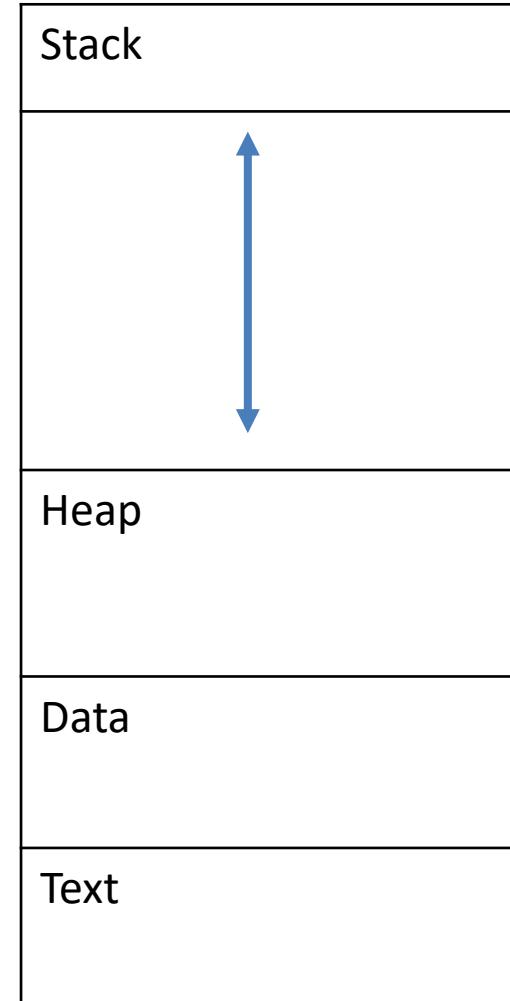
# Are we secure now?

---

- Confinement problem [3]
  - Deals with preventing a process from transmitting information to any other program except its caller.
- Covert channels
  - Path of communication that was not designed to be used for communication.

# Processes

- A program in execution
- **Stack:** temporal data e.g. local variables, return address
- **Heap:** dynamic allocated memory of the process at run time
- **Text:** current activity by the program counter and registers
- **Data:** contains global and static variables



# Threads

- A process generates a thread when it requires to send something to the CPU for processing
- The thread includes an individual instruction set and data
- Generated dynamically
- Multithreaded applications are capable of running several different threads at the same time.
- Threads share the same resources of the process that created them

# Memory leaks

---

- Every process is allocated with an amount of memory
- Memory should be released when the process finished with it *Responsibility of the user*
- Not the case for poorly written applications
- May be used for Denial of Service attacks and lead to memory starvation



# Technical attacks

---

- Buffer overflow attacks
- String formatting attacks
- Integer overflow attacks
- Code injection attacks

# Buffer overflow

- C compiler does not do array bounds checking
- Suppose the following code

```
int i;  
  
char ch[16];  
  
i=32;  
  
ch[i]=0;
```

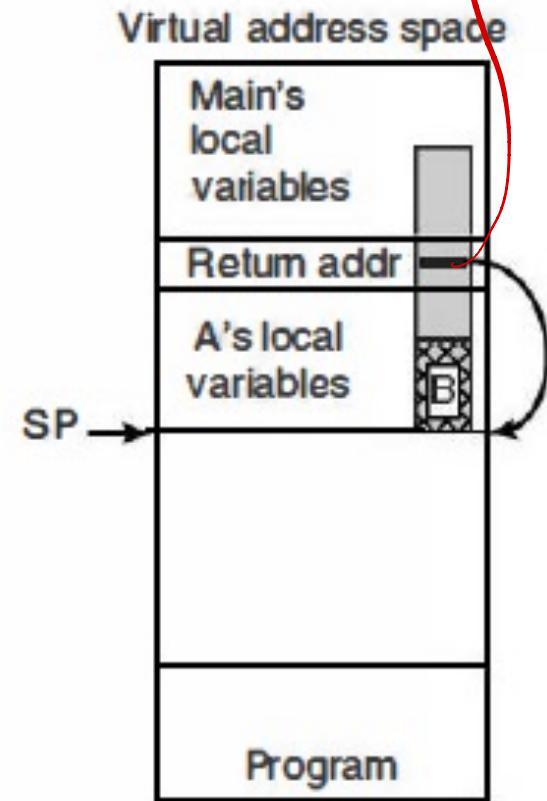
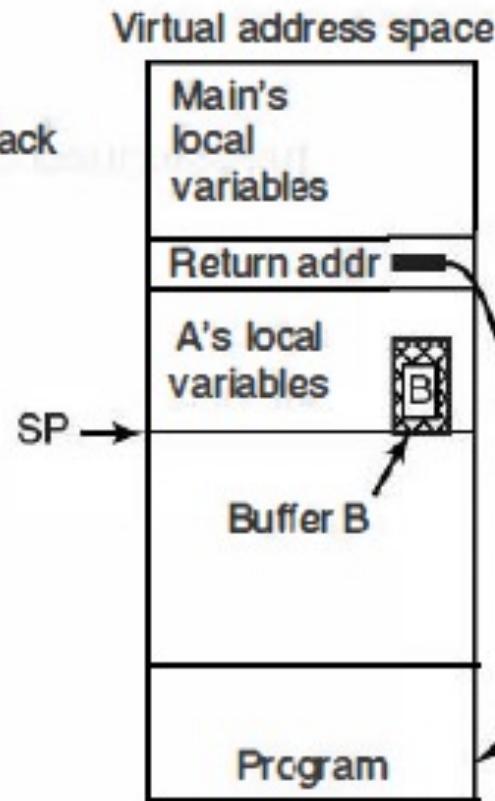
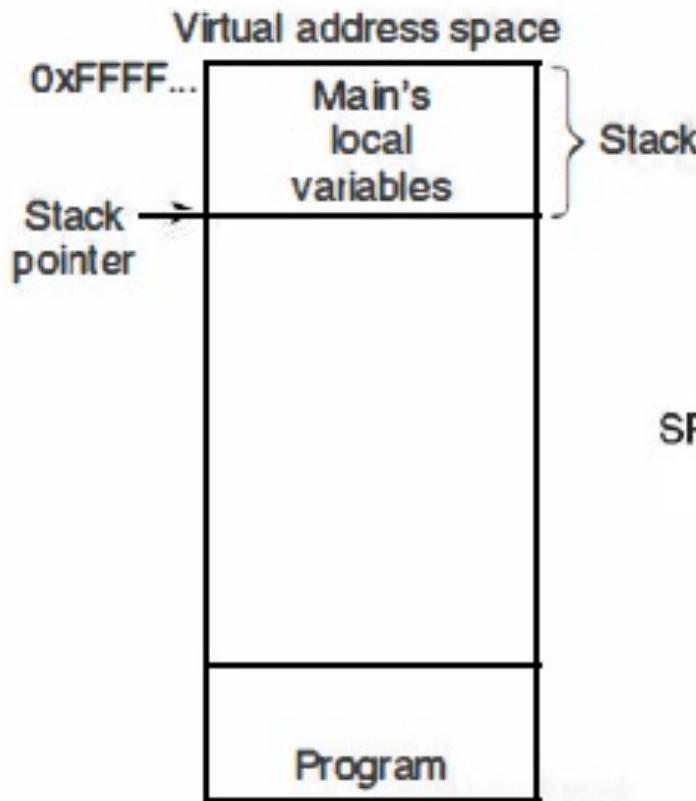
Possible Segmentation fault



# Buffer overflow

Oversizing return address  
segmentation fault

[source: 1]



# String formatting attacks

---

- Issues when valid formatting is not used with `printf`
- Statements such as `printf(buffer)` are still valid
- An attacker can pass a list of parameters that can be executed as a command
  - Execute code, read the stack, etc.

# String formatting attacks

```
#include <stdio.h>
int main() {
    char* s = "%x %x Hello World";
    printf("%s\n", s); // %x %x Hello World
    printf(s); //aedda000 aebb49e0 Hello World
    return 0;
}
```

# Integer overflow

- Integer arithmetic operations are commonly done using modulo arithmetic
- Modulo arithmetic allows values to wrap if they go above a certain value
- 8-bit integer will hold values between 0 and 255
- What if a higher number has to be stored there?

Depends on the language  
hardware may<sup>\*</sup> raise error

# Code injection attacks

- Gets a program to execute code without realising that
- Mainly due to poor implementations

```
int main(void) {  
    char command[1024], src[500], dst[500];  
    strcpy(command, "cp ");  
    printf("Source: "); gets(src);  
    printf("Destination: "); gets(dst);  
    strcpy(command, src); strcpy(command, " ");  
    strcpy(command, dst);  
    system(command);  
    return 0; }
```

# Code injection attacks

---

- If a user adds
  - src = source.txt and dst = destination.txt
- It will execute
  - cp source.txt destination.txt
- How about
  - src = source.txt and
  - dst = destination.txt; rm -rf /
- It will execute
  - cp source.txt destination.txt; rm -rf /

# Bounded Software Model Checking

## An Introduction to CBMC

**SCC.363 Security and Risk**

*School of Computing and Communications, Lancaster University*

14th February 2023

# Boolean Satisfiability (SAT) Problem

The Boolean satisfiability problem (also **SAT problem**) asks:

*Is a given Boolean formula  $F$  satisfiable?*

E.g. consider the following Boolean formula:

$$F = (x_1 \wedge x_2 \vee x_3) \wedge (x_4 \wedge x_1 \vee x_3) \wedge \neg(x_4 \vee x_1)$$

Problem: does there exist an assignment of truth values to the Boolean variables  $x_1, x_2, x_3, x_4$  that makes the formula True?

**Yes!**  $x_1 = \text{False}$ ,  $x_2 = \text{True}$ ,  $x_3 = \text{True}$ ,  $x_4 = \text{False}$ .

So  $F$  is *satisfiable* (otherwise *unsatisfiable*).

# Boolean Satisfiability (SAT) Problem

The Boolean satisfiability problem (also **SAT problem**) asks:

*Is a given Boolean formula  $F$  satisfiable?*

E.g. consider the following Boolean formula:

$$F = (x_1 \wedge x_2 \vee x_3) \wedge (x_4 \wedge x_1 \vee x_3) \wedge \neg(x_4 \vee x_1)$$

Problem: does there exist an assignment of truth values to the Boolean variables  $x_1, x_2, x_3, x_4$  that makes the formula True?

**Yes!**  $x_1 = \text{False}$ ,  $x_2 = \text{True}$ ,  $x_3 = \text{True}$ ,  $x_4 = \text{False}$ .

So  $F$  is *satisfiable* (otherwise *unsatisfiable*).

The SAT problem was the first example of an NP-Complete problem (Cook-Levin theorem, 1971).

No known algorithm can solve it in polynomial time

There are some SAT solvers

# Boolean Satisfiability (SAT) Problem

Many problems can be encoded in terms of Boolean satisfiability.

E.g. many problems in circuit design can be reduced to SAT.

An output from a wire in a digital circuit corresponds to a Boolean formula whose variables correspond to the input wires (along the wires 1 = True, 0 = False).

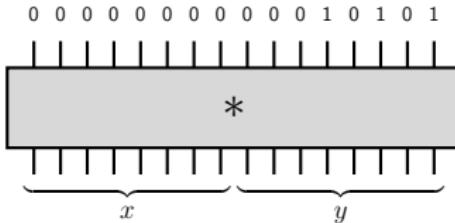
# Boolean Satisfiability (SAT) Problem

Many problems can be encoded in terms of Boolean satisfiability.

E.g. many problems in circuit design can be reduced to SAT.

An output from a wire in a digital circuit corresponds to a Boolean formula whose variables correspond to the input wires (along the wires 1 = True, 0 = False).

One may e.g. model a *multiplier circuit* as a Boolean formula and ask if there are inputs to the circuit that result in the number 21 (10101 in binary) along the output wires (i.e. check if 21 is *prime*).



# SAT Solvers

Boolean formula  $F \longrightarrow$  **SAT solver**  $\longrightarrow$  SAT/UNSAT

Modern SAT solvers are very efficient and can solve SAT problems with *millions* of Boolean variables! Enormous progress made in the past 25 years, spurred on by the SAT competition.

<http://www.satcompetition.org/>

Most modern SAT solvers are based on the **DPLL** algorithm or its refinements such as **CDCL**, with exponential worst-case time complexity.

**DPLL** is based on *backtracking* and works on Boolean formulas in *Conjunctive Normal Form* (CNF).

# Conjunctive Normal Form

A literal is a Boolean variable  $x_i$ , or its negation  $\neg x_i$ .

A clause is a disjunction of literals (e.g.  $x_1 \vee \neg x_3 \vee x_2$ ).

A Boolean formula is in Conjunctive Normal Form (CNF) if it is a conjunction of one or more clauses, e.g.

$$(x_1 \vee \neg x_3 \vee x_2) \wedge (\neg x_4 \vee \neg x_1 \vee x_3) \wedge (\neg x_5).$$

Any Boolean formula  $F$  can be converted into CNF by applying some simple transformations:

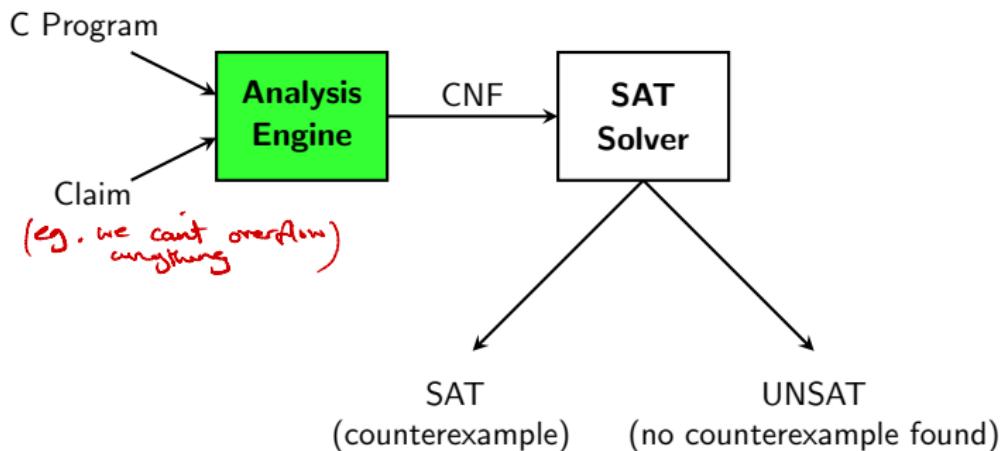
- eliminating double negations (i.e.  $\neg\neg P$  becomes  $P$ ),
- De Morgan's laws (e.g.  $\neg(P \wedge Q) = \neg P \vee \neg Q$ ),
- Distributive laws (e.g.  $P \vee (Q \wedge R) = (P \vee Q) \wedge (P \vee R)$ , etc.)

# An Application of SAT Solvers: Model Checking

**CBMC**: the C Bounded Model Checker

(<https://www.cprover.org/cbmc/>)

Main idea: Given a C program and a claim, use a SAT solver to check if there is an execution that violates the claim.



# CBMC: Bug [Catching 😊/ Finding 🚫] with a SAT Solver

Developed by D. Kroening and others at CMU in 2003.

Given a C program, CBMC can automatically check simple safety claims, some of which are important to security:

- array bound checks,
- division by zero,
- arithmetic overflow,
- pointer checks (NULL pointer dereference),
- user-supplied assertions.

CBMC expects there to be a program entry point, i.e. a `main()`.

It allows the user to make *assertions* using `assert()` and to create *assumptions* using `_CPROVER_assume()`.

# Using CBMC

Claims made by decorating code with assumptions and assertions.

- **assert(e)**

*aborts execution when e is false; no-op otherwise.*

```
void assert(_Bool e) { if (!e) exit(); }
```

- **\_CPROVER\_assume(e)**

*"ignores" execution when e is false; no-op otherwise.*

Program traces are restricted to those satisfying the assumption.

To find *counterexamples* to claims in a program `prog.c` we run:

```
$ cbmc --trace prog.c
```

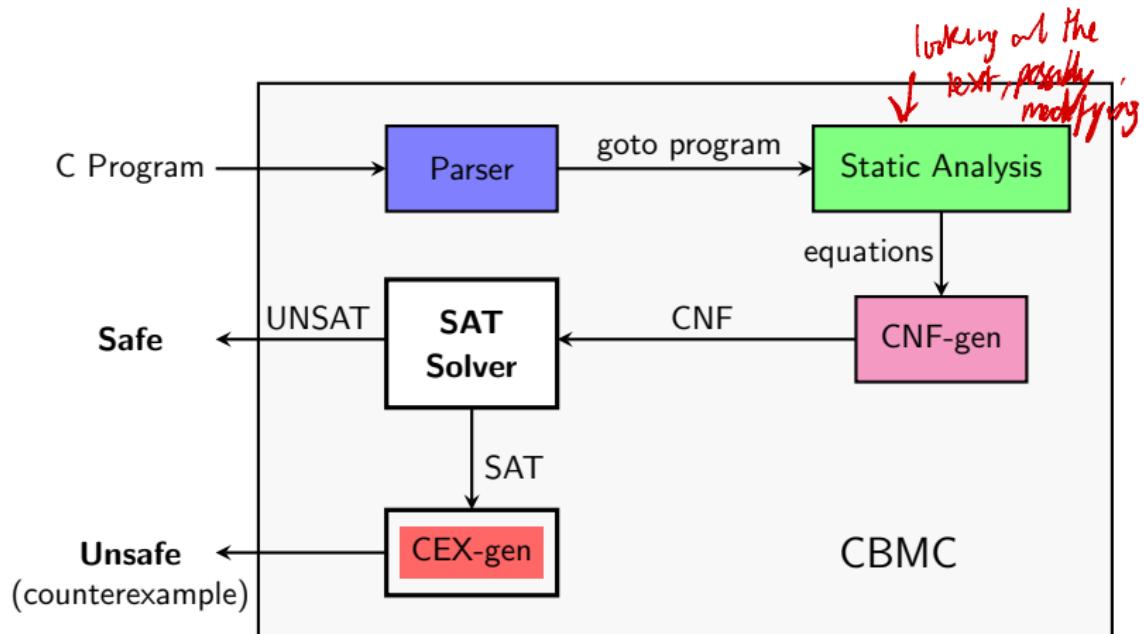
# CBMC: How Does It Work?

CBMC works by transforming a C program into a set of equations.

The main steps in CBMC are:

- 1 Simplify control flow
- 2 Unwind all the loops
- 3 Convert into *Single Static Assignment (SSA)*
- 4 Convert into equations
- 5 “Bit-blast”
- 6 Solve with a SAT solver
- 7 Convert SAT assignment into a counterexample

# CBMC: How Does It Work?



# Control Flow Simplification

- All side effects are removed

e.g.  $j=i++;$  is transformed into  $j=i;$   $i=i+1;$

- Control flow is made *explicit*

**continue** and **break** are replaced by **goto**

- All loops are simplified into *one form*

e.g. **for**, **do**, **while** are replaced by just **while**

# Loop Unwinding

- All loops are *unwound*:
  - *can use different unwinding bounds for different loops,*
  - *can check whether unwinding is sufficient using a special unwinding assertion.*
- If a program satisfies all of its claims *and* all unwinding assertions, then it is *correct*.
- Recursive functions and backward **goto** are similar (use *inlining*).

# Loop Unwinding

**while** loops are unwound *iteratively*.

(**break/continue** replaced by **goto**.)

```
1 void f(...){  
2     ... // some code  
3     while(cond){  
4         Body;  
5     }  
6     Remainder;  
7 }
```

```
1 void f(...){  
2     ... // some code  
3     if(cond){  
4         Body;  
5         if(cond){  
6             Body;  
7             assert(!cond);  
8         } //Unwinding assertion  
9     }  
10    }  
11    }  
12    }  
13    Remainder;  
14 }
```

Assertion inserted after last iteration: violated if the program runs longer than bound permits.

# Loop Unwinding

Assertion inserted after last iteration: violated if the program runs longer than bound permits. Positive correctness result!

```
1 void f(...){  
2     ... // some code  
3     if(cond){  
4         Body;  
5         if(cond){  
6             Body;  
7             if(cond){  
8                 Body;  
9                 assert(!cond); //Unwinding assertion  
10            }  
11        }  
12    }  
13    Remainder;  
14 }
```

## Example: Sufficient Loop Unwinding

unwind = 3

```
1 void f(...){  
2     j = 1;  
3     while(j<=2){  
4         j = j + 1;  
5     }  
6     Remainder;  
7 }
```

```
1 void f(...){  
2     j = 1;  
3     if(j<=2){  
4         j = j + 1;  
5         if(j<=2){  
6             j = j + 1;  
7             if(j<=2){  
8                 j = j + 1;  
9                 assert(!(j<=2));  
10            }  
11        }  
12    }  
13    Remainder;  
14 }
```

## Example: Insufficient Loop Unwinding

unwind = 3

```
1 void f(...){  
2     j = 1;  
3     while(j<=10){  
4         j = j + 1;  
5     }  
6     Remainder;  
7 }
```

no. If statements

```
1 void f(...){  
2     j = 1;  
3     if(j<=10){  
4         j = j + 1;  
5         if(j<=10){  
6             j = j + 1;  
7             if(j<=10){  
8                 j = j + 1;  
9                 assert(!(j<=10));  
10            }  
11        }  
12    }  
13    Remainder;  
14 }
```

# Transforming Loop-Free Programs Into Equations

It is trivial to translate a program into a set of equations if each variable is only assigned once!

```
1  x = a;  
2  y = x+1;  
3  z = y-1;
```

This program is directly transformed into

$$x = a \wedge y = x + 1 \wedge z = y - 1.$$

# Transforming Loop-Free Programs Into Equations

## Static Single Assignment (SSA) form.

- Every variable is assigned exactly once.
- Every variable is defined before it is used.

When a variable is assigned multiple times, we use a new variable for each assignment.

```
1  x=x+y;  
2  x=x*2;  
3  a[i]=100;
```

```
1  x1 = x0 + y0;  
2  x2 = x1*2;  
3  a1[i0] = 100;
```

What about conditionals?

# Transforming Loop-Free Programs Into Equations

Converting conditionals to SSA.

```
1  if (v)
2    x = y;
3  else
4    x = z;
5  w = x;
```

```
1  if (v0)
2    x0 = y0;
3  else
4    x1 = z0;
5  w1 = x?? // which x?
```

# Transforming Loop-Free Programs Into Equations

Converting conditionals to SSA.

```
1  if (v)
2    x = y;
3  else
4    x = z;
5  w = x;
```

```
1  if (v0)
2  x0 = y0;
3  else
4  x1 = z0;
5  x2 = v0 ? x0 : x1;
6  w1 = x2;
```

For each joint point add new variables with *selectors*.

# Loop-Free Example

Starting from the following C code:

```
1 int y;
2 int x;
3 x=x+y;
4 if(x != 1)
5     x=2;
6 else
7     x++;
8 assert(x<=3);
```

# Loop-Free Example

Simplify control flow and remove side effects

```
1 int y;
2 int x;
3 x=x+y;
4 if(x != 1)
5     x=2;
6 else
7     x=x+1;
8 assert(x<=3);
```

Convert to SSA (Static Single Assignment form)

```
1 x1 = x0+y0;
2 if(x1 != 1)
3     x2 = 2;
4 else
5     x3 = x1 + 1;
6 x4 = (x1 != 1) ? x2 : x3;
7 assert(x4<=3);
```

# Loop-Free Example

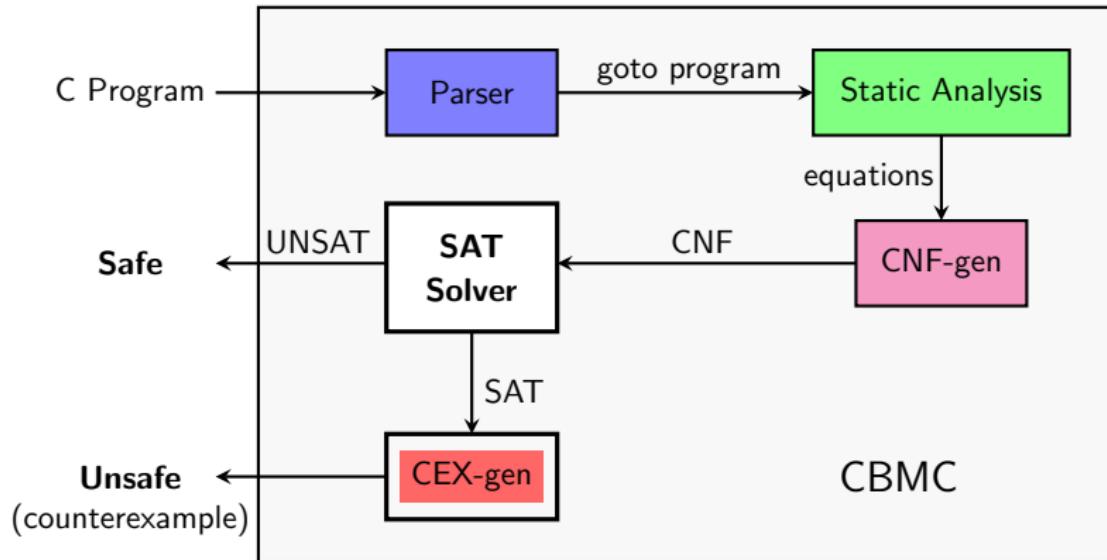
Convert to SSA (Static Single Assignment form)

```
1  x1 = x0+y0;
2  if(x1 != 1)
3      x2 = 2;
4  else
5      x3 = x1 + 1;
6  x4 = (x1 != 1) ? x2 : x3;
7  assert(x4<=3);
```

Generate constraints (if SAT, then assertion is false):

$$\begin{aligned} &x_1 = x_0 + y_0 \wedge x_2 = 2 \wedge x_3 = x_1 + 1 \\ &\wedge ((x_1 \neq 1 \wedge x_4 = x_2) \vee (x_1 = 1 \wedge x_4 = x_3)) \quad [\text{selector}] \\ &\wedge \neg(x_4 \leq 3) \quad [\text{negated assertion}] \end{aligned}$$

# CBMC: How Does It Work?



## Bit Blasting

So far, formulas such as  $x_2 = x_1 + 1 \wedge y_2 = x_2$  are *not* stated in propositional logic!

The operations are performed on *bit vectors*.

In order to convert these formulas into a format acceptable to a SAT solver, one needs to apply *flattening/bit blasting*.

# Bit Blasting

So far, formulas such as  $x_2 = x_1 + 1 \wedge y_2 = x_2$  are *not* stated in propositional logic!

The operations are performed on *bit vectors*.

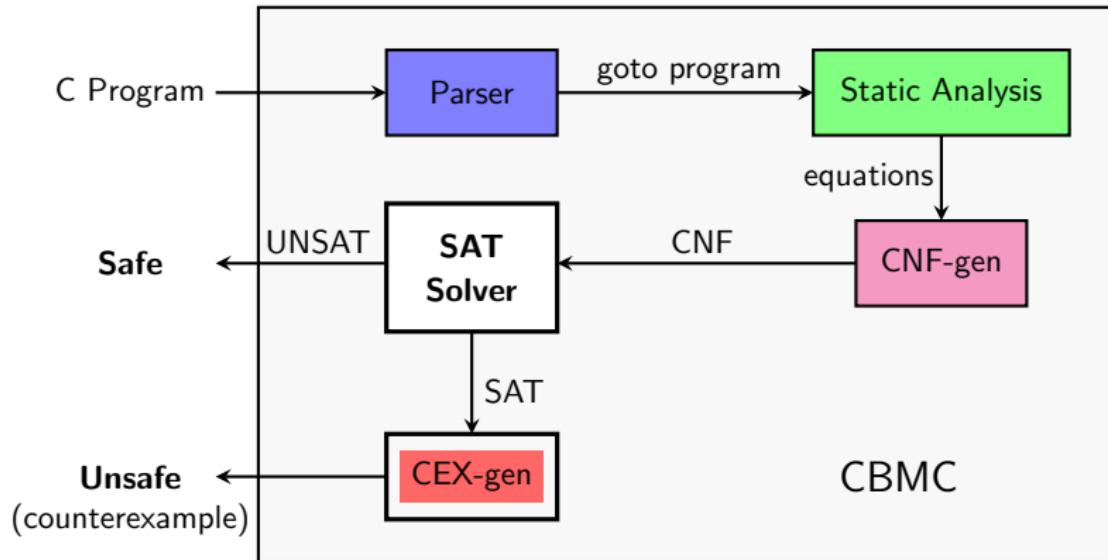
In order to convert these formulas into a format acceptable to a SAT solver, one needs to apply *flattening/bit blasting*.

Intuitively, we can build Boolean circuits for the bit-vector operations; these can be described by Boolean formulas.

*Unfortunately with a lot more variables.*

- breaking the formulas down to their smallest parts (bits) & representing them using simple operations like AND, OR, and NOT
- computers can effectively evaluate the formulas, check if they're true/false

# CBMC: How Does It Work?

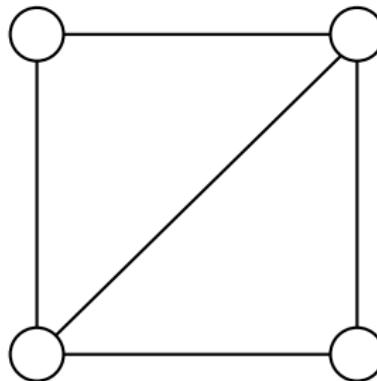


# Challenge Problem: Graph Colouring using CBMC

## Graph colouring problem:

Given a finite undirected graph  $G$  and a finite set of colours  $C$ , how can we use CBMC to solve the colouring problem, i.e. to check if there is an assignment of colours from  $C$  to nodes in the graph  $G$  such that no two connected nodes have the same colour?

E.g. let us consider the following  $G$ , with  $C = \{\text{red}, \text{green}, \text{blue}\}$ .



## Further Reading:

- CBMC Tutorial:  
<http://www.cprover.org/cprover-manual/cbmc/tutorial/>
- **Edmund Clarke, et al.** "*Behavioral consistency of C and Verilog programs using bounded model checking.*" " Proceedings 2003. Design Automation Conference. IEEE, 2003.  
[CMU-CS-03-126.pdf](#)

### Optional :

- **Lucas Cordeiro, et al.** "*SMT-based bounded model checking for embedded ANSI-C software.*" IEEE Transactions on Software Engineering 38.4 (2011): 957-974.  
<https://core.ac.uk/download/pdf/59348834.pdf>

*Acknowledgements:* Partly based on material by D. Kroening and G. Parlato.